Out-of-Distribution Learning

Gaussian Tasks Experiment

- Consider an in-distribution task that consists of two class conditional gaussians.
- Now, consider an out-of-distribution task similar to the above task, but whose center is displaced by an amount Δ.
- The amount Δ reflects the "similarity" between the two tasks.

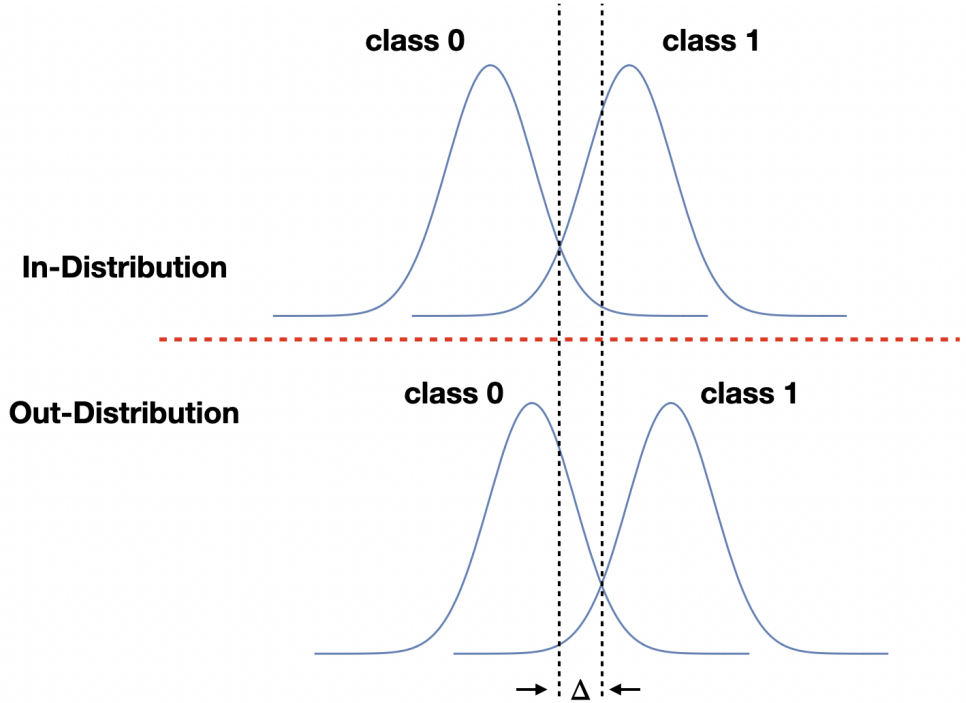

Gaussian Tasks Experiment

- We have access to n samples from the in-distribution task, and m samples from the out-of-distribution task.
- Using both the in-distribution and out-of-distribution samples, we train a classifier h aimed at the in-distribution classification task.
- Let's denote the classification error of h by $\mathbb{E}[L_{n,m,\Delta}]$.

Gaussian Tasks Experiment

- Let n be a small fixed constant. We hypothesize that,
  - For very small Δ, as we add more out-of-distribution data (as m increases) the $\mathbb{E}[L_{n,m,\Delta}]$ would decrease.
  - For moderately large Δ, as we add more out-of-distribution data (as m increases) the $\mathbb{E}[L_{n,m,\Delta}]$ would initially decrease and start increasing later. The initial decrease is due to the reduction in the variance of h. The later increase is due to the increase in bias of h caused by the out-of-distribution samples.
  - For very large Δ, as we add more out-of-distribution data (as m increases) the $\mathbb{E}[L_{n,m,\Delta}]$ would keep increasing.

Gaussian Tasks Experiment

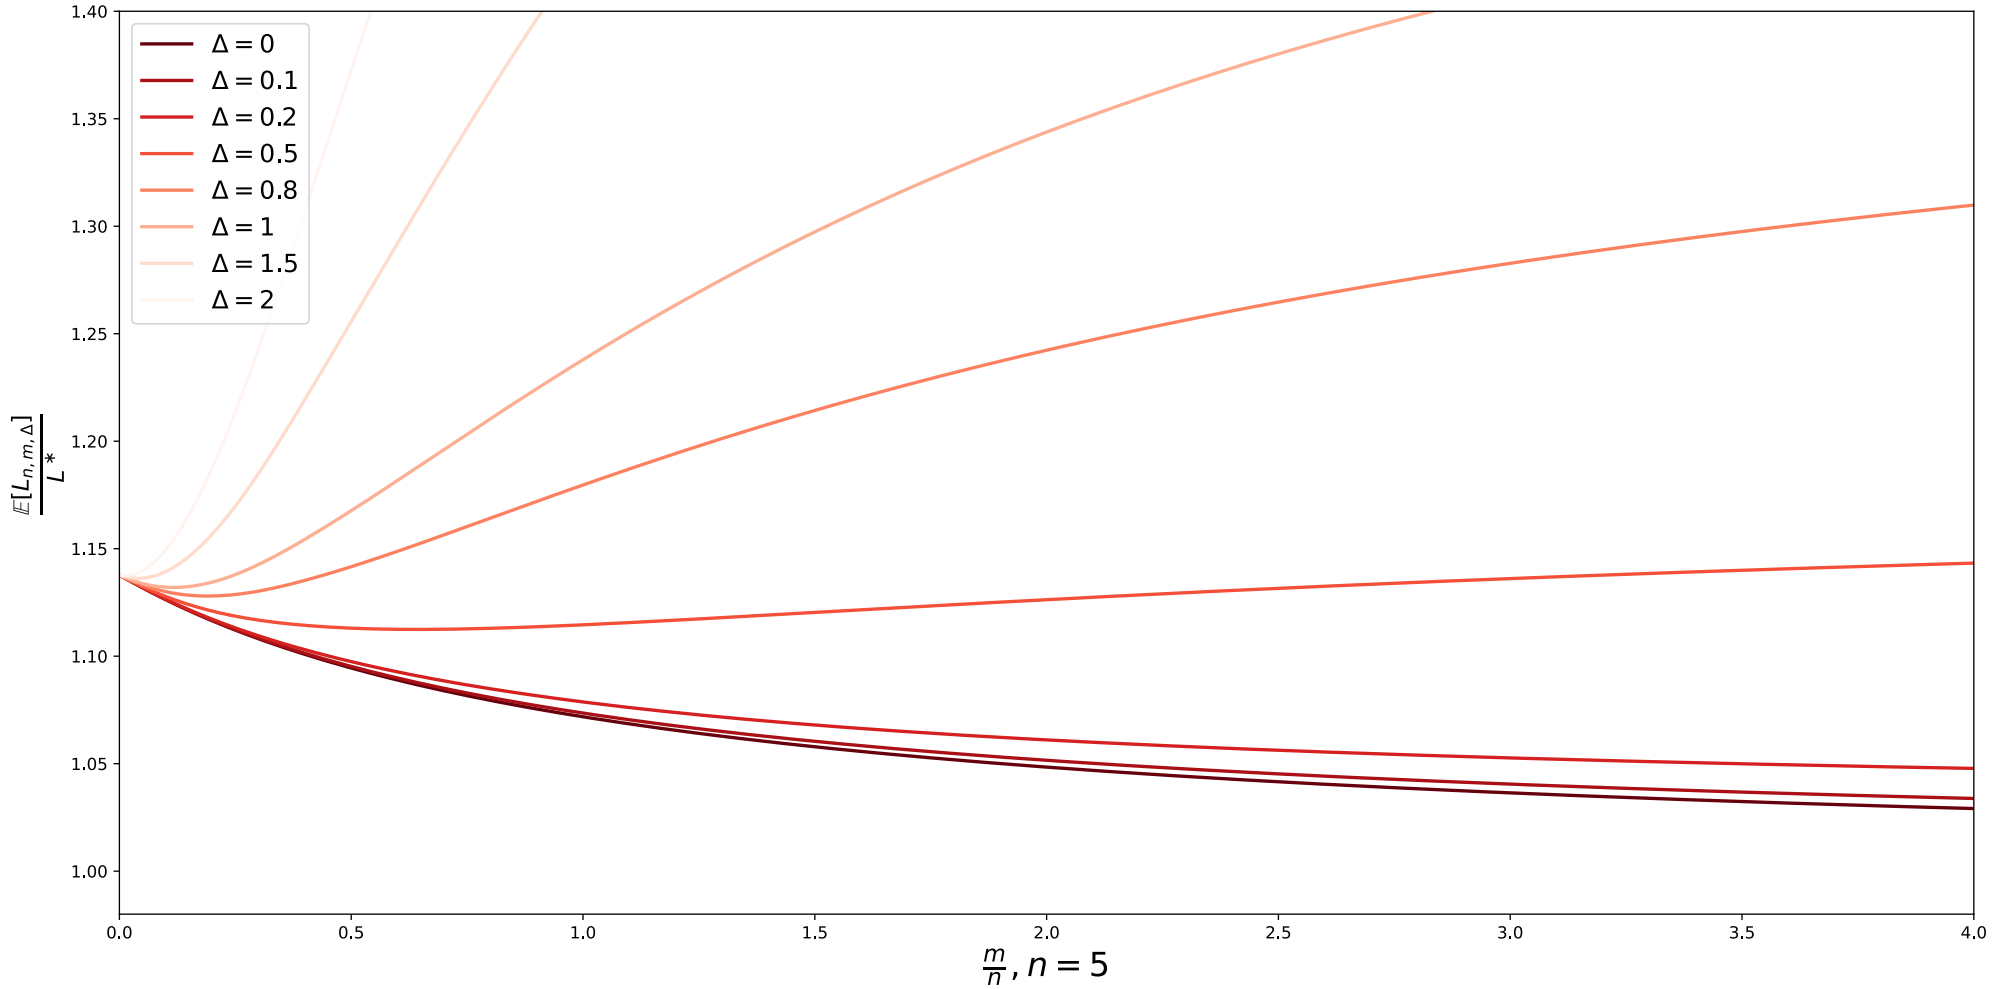

Gaussian Tasks Experiment

- Number of replicates: 1000

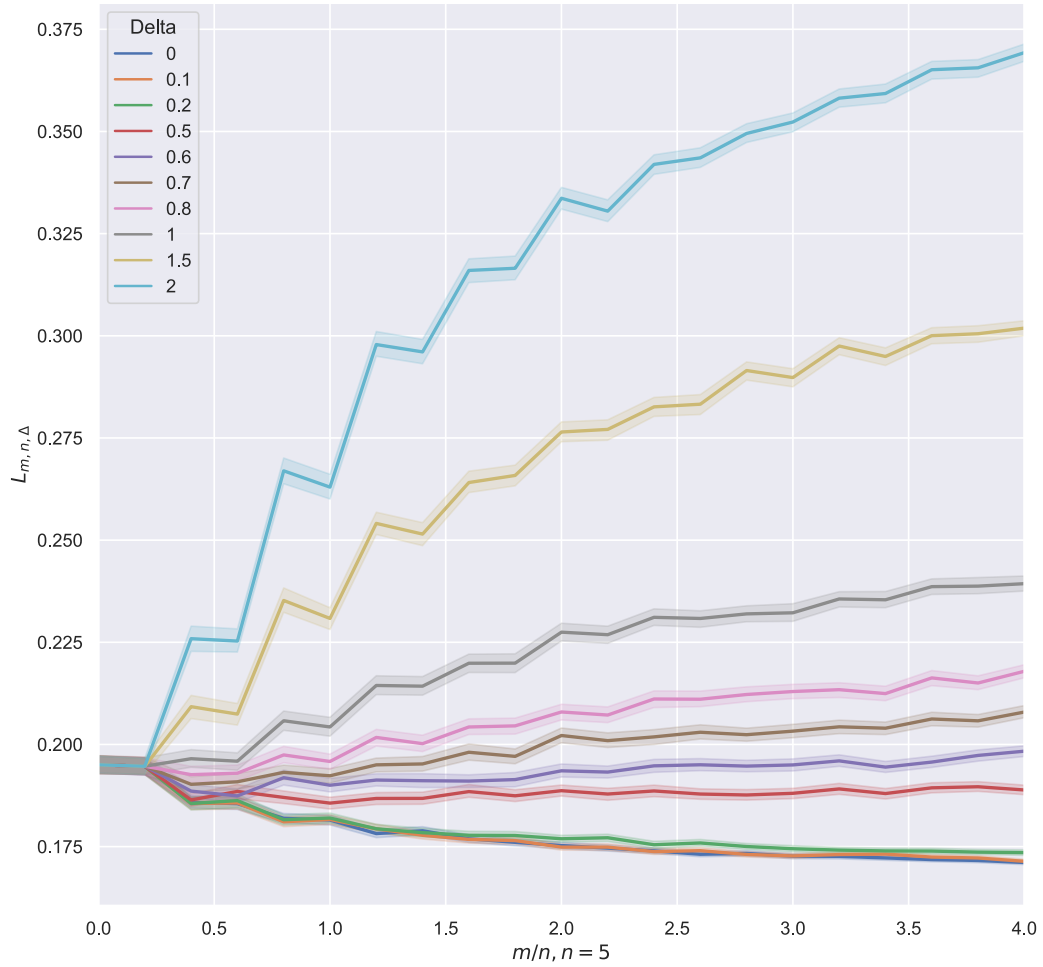

Gaussian Tasks Experiment

- Number of replicates: 1000

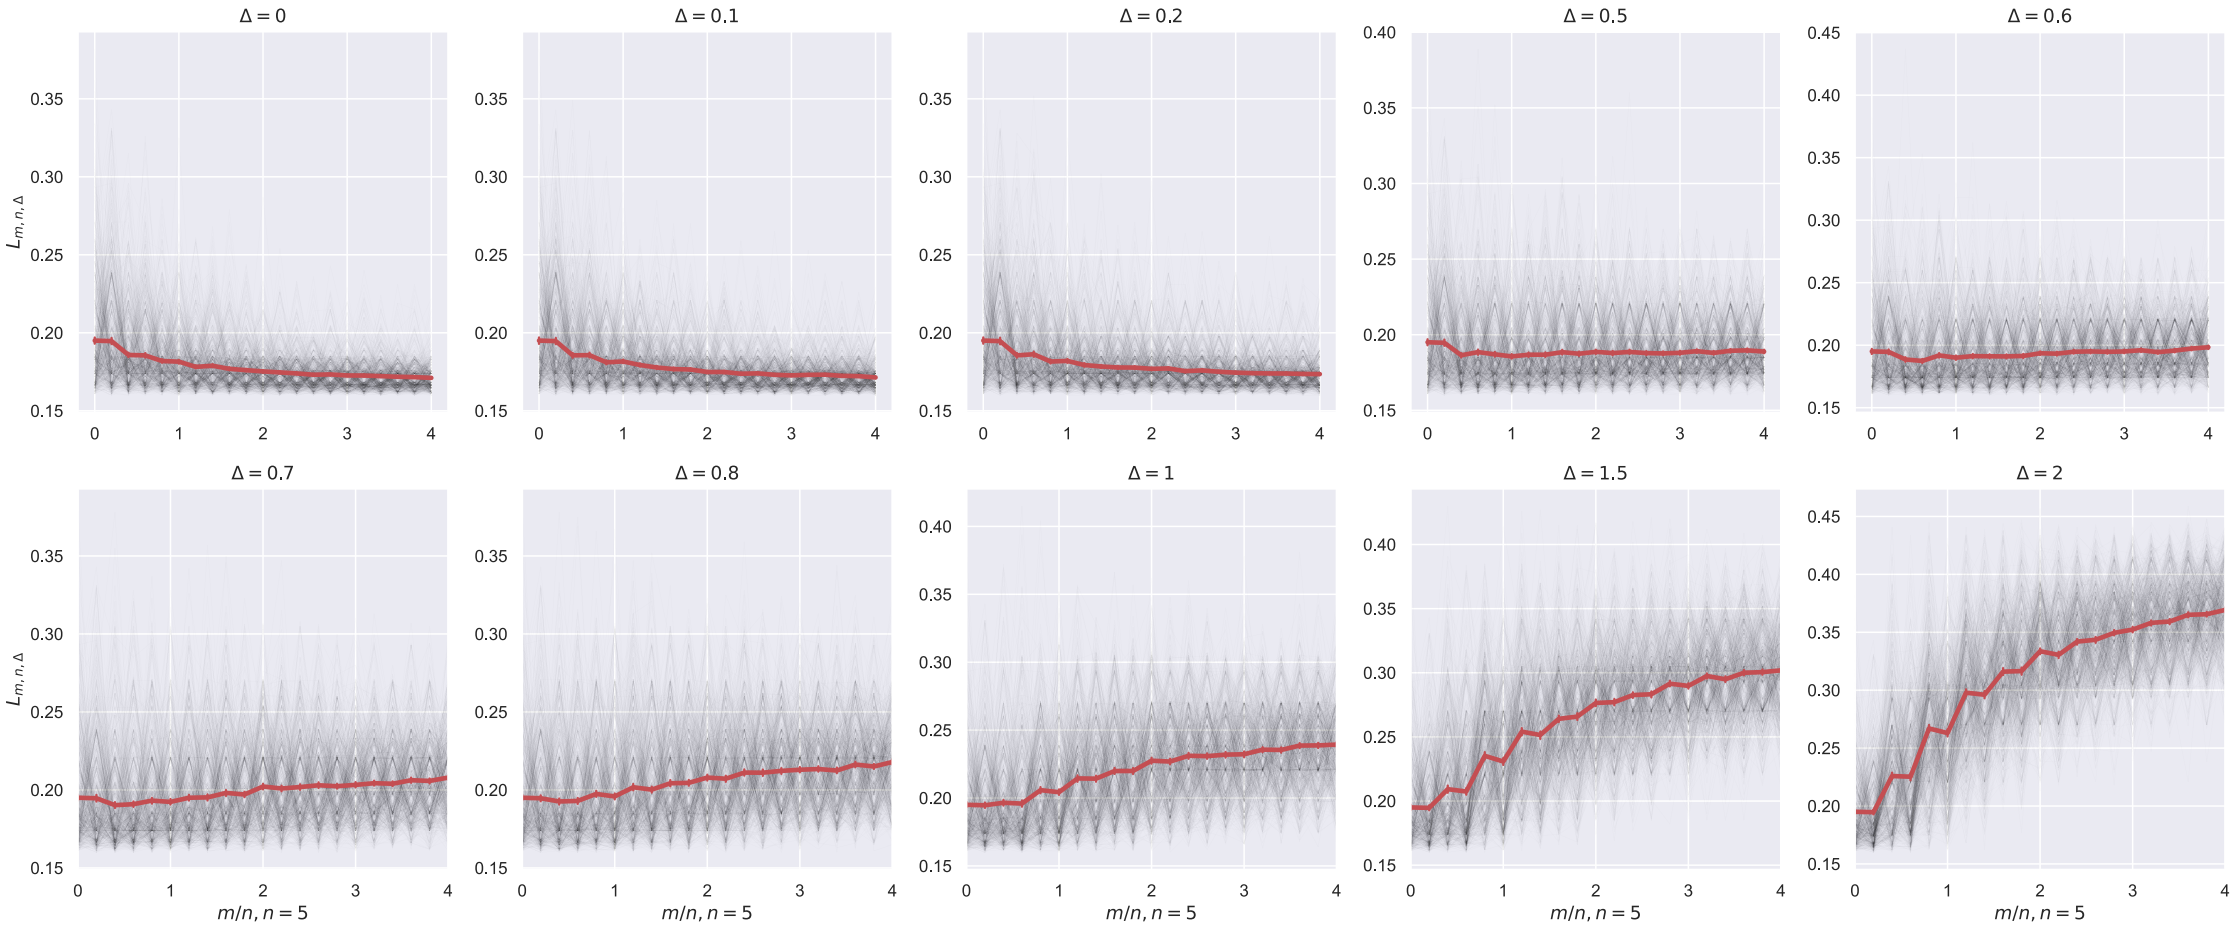

Bird vs. Cat & α-Rotated Bird vs. Cat (Single-Head Network)

- Number of replicates: 10, Network: SmallConv

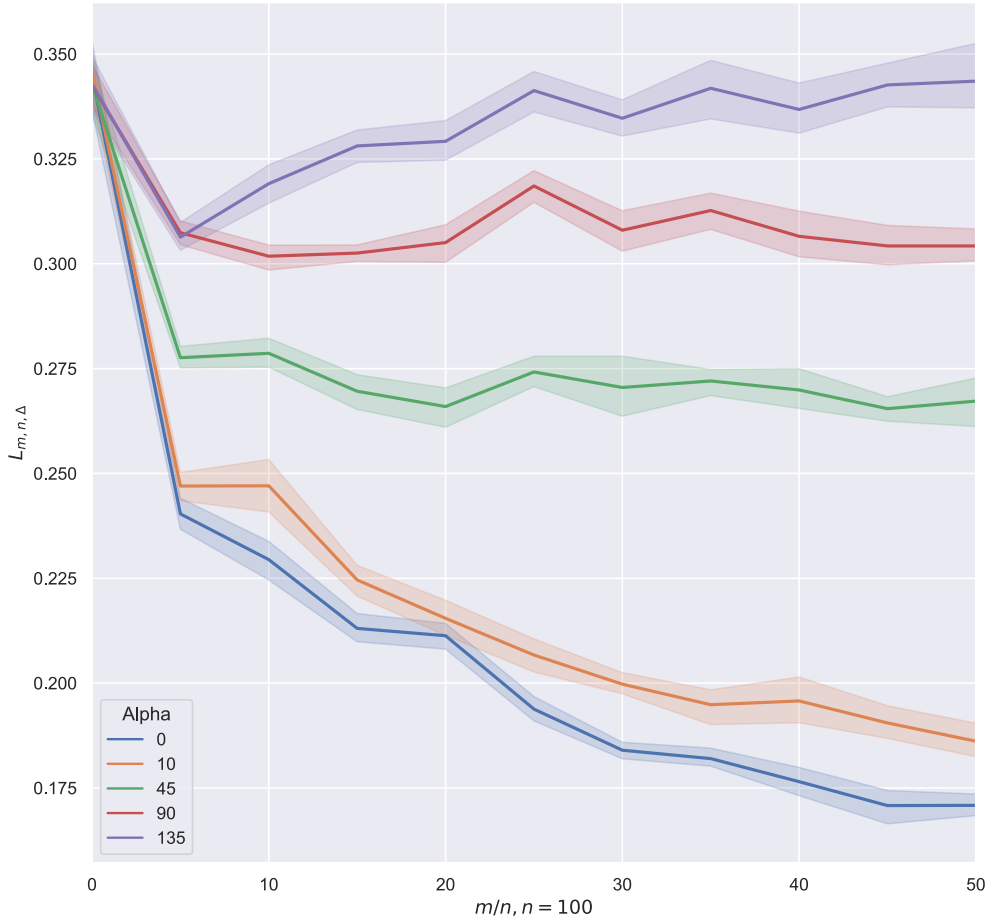

Bird vs. Cat & α-Rotated Bird vs. Cat (Single-Head Network)

- Number of replicates: 10, Network: SmallConv

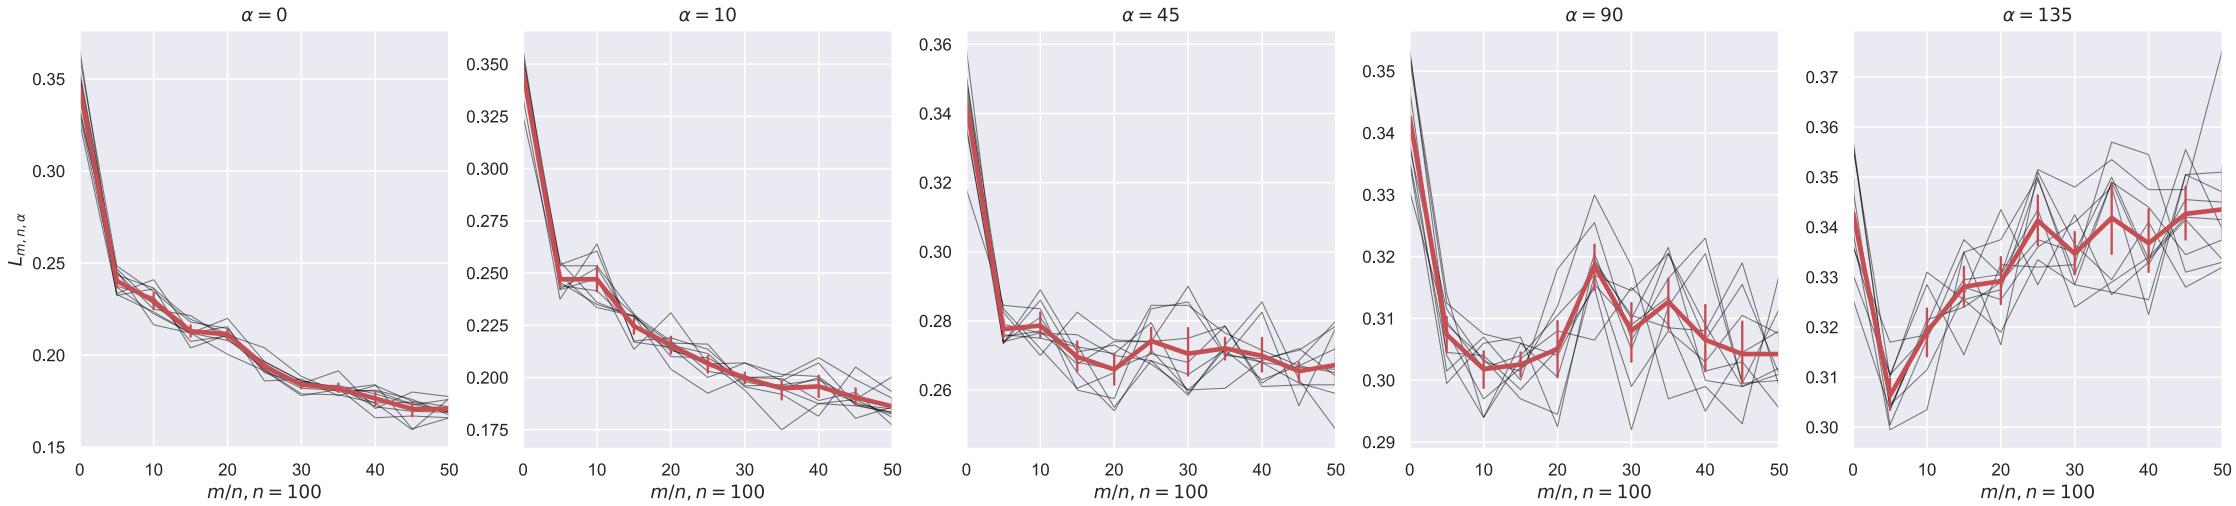

Task 2: Bird vs. Cat & Task 3: Deer vs. Dog (Single-Head Network)

- Number of replicates: 20, Network: SmallConv

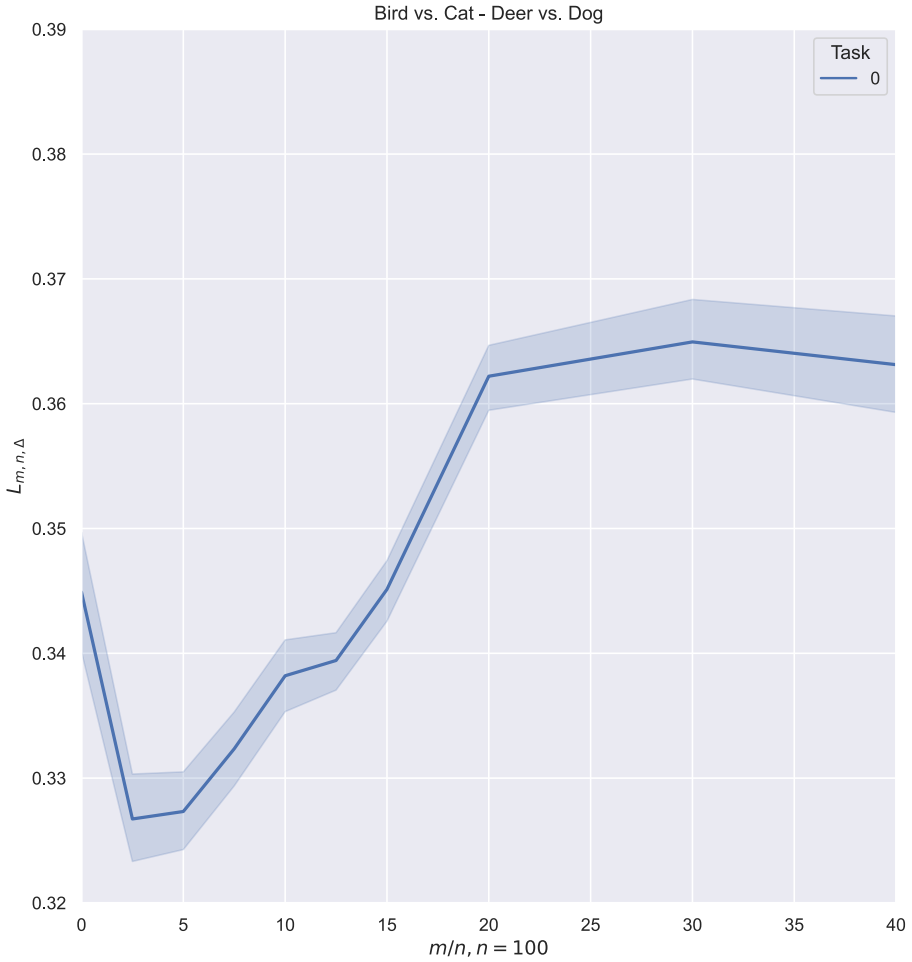

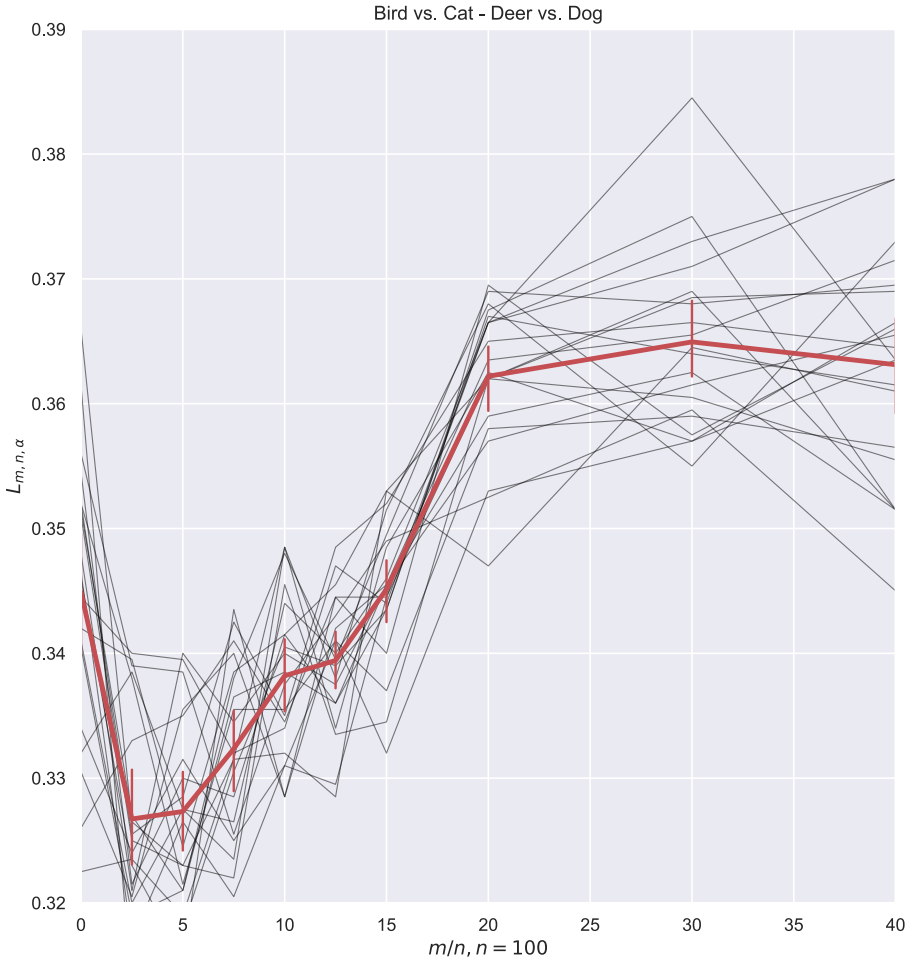

Task 2: Bird vs. Cat & Task 3: Deer vs. Dog (Single-Head Network)

- Number of replicates: 20, Network: SmallConv, each model was trained for 100 epochs

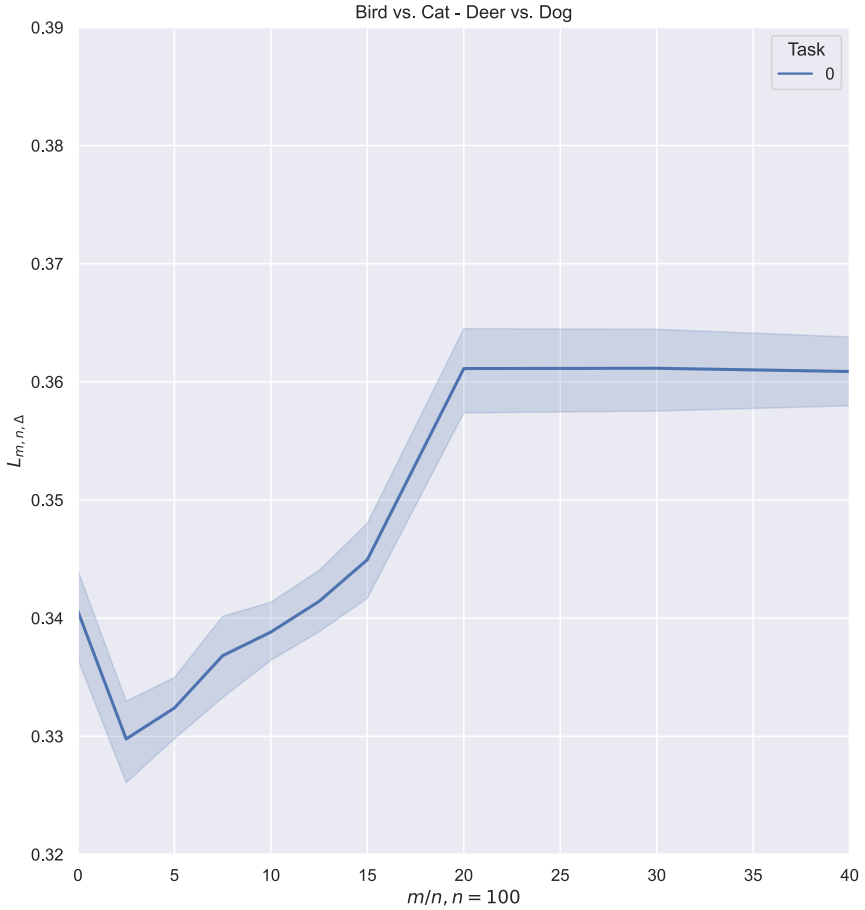

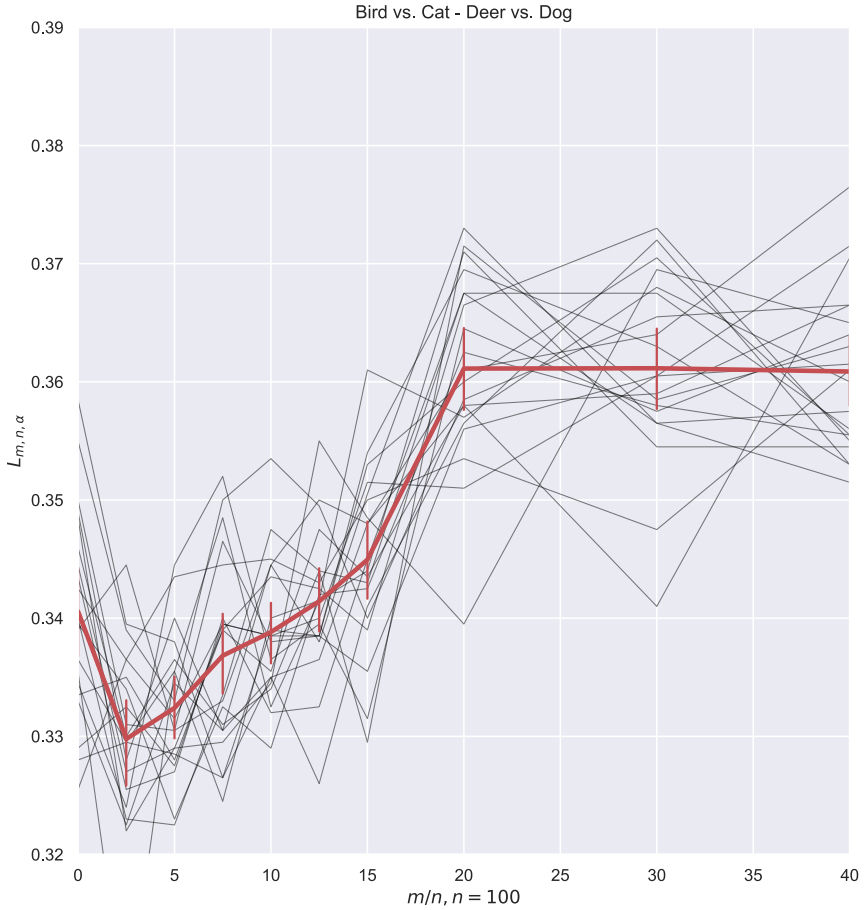

Task 2: Bird vs. Cat & Task 3: Deer vs. Dog (Multi-Head Network)

- Number of replicates: 20, Network: SmallConv

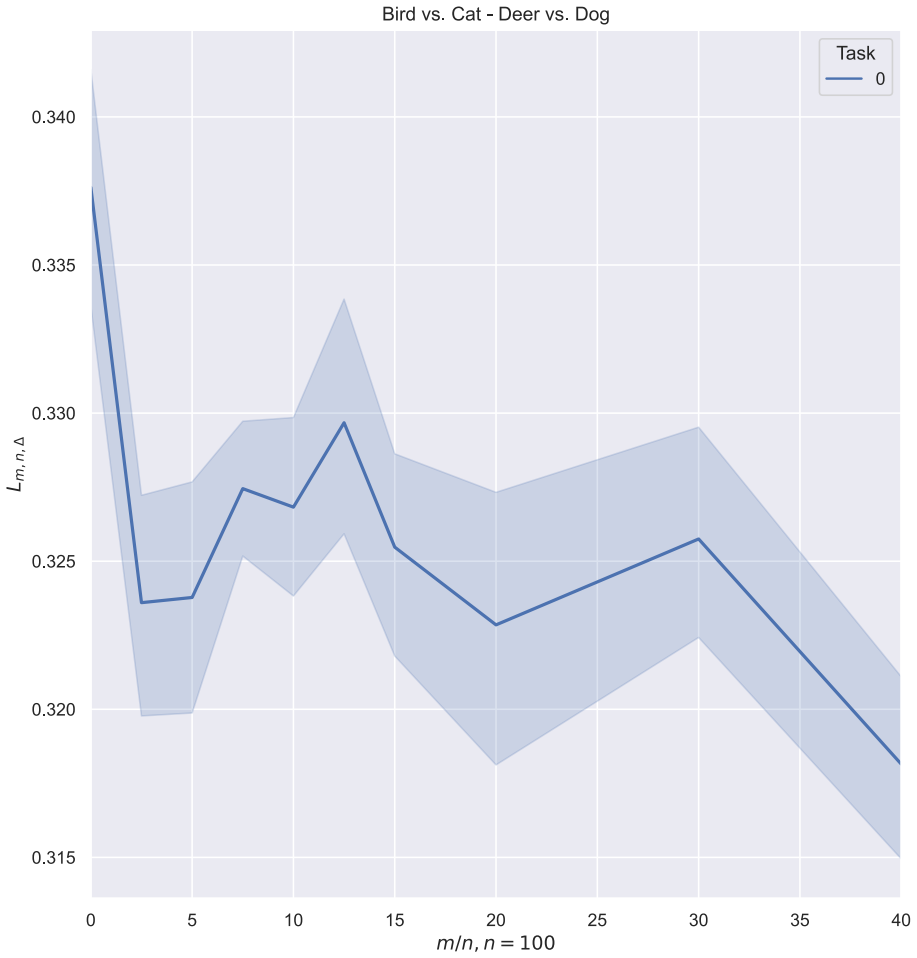

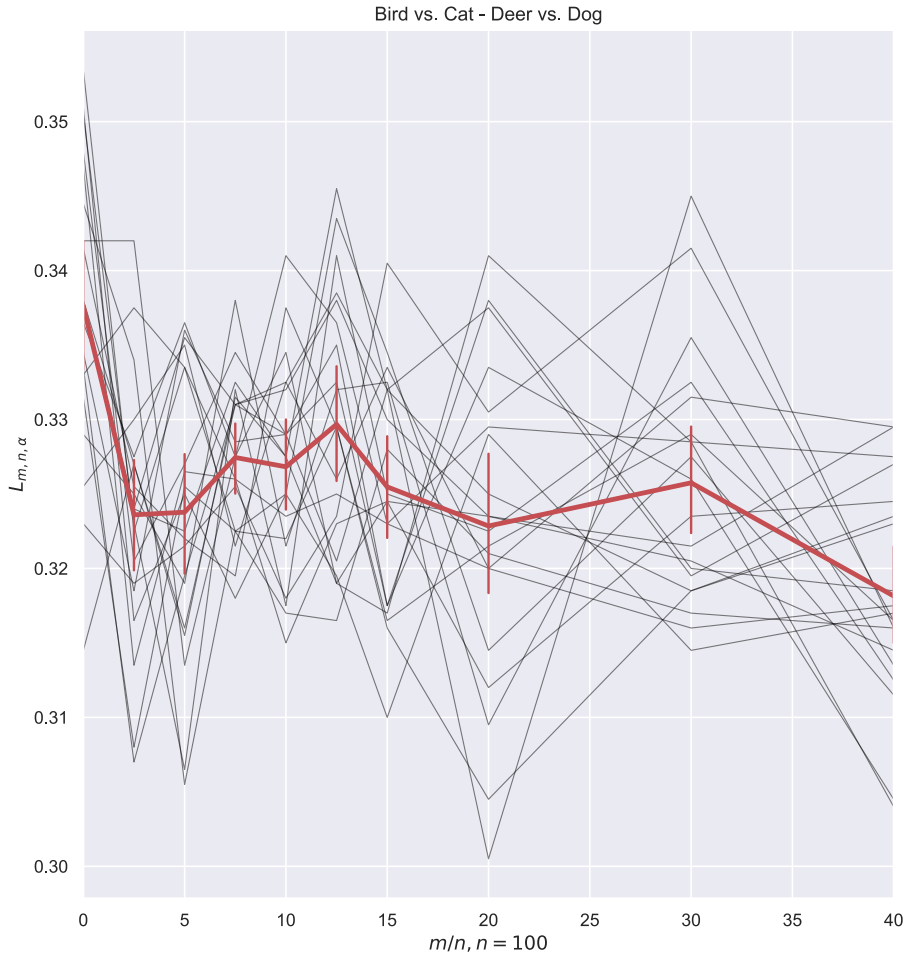

Task 2: Bird vs. Cat & Task 4: Frog vs. Horse (Multi-Head Network)

- Number of replicates: 20, Network: SmallConv

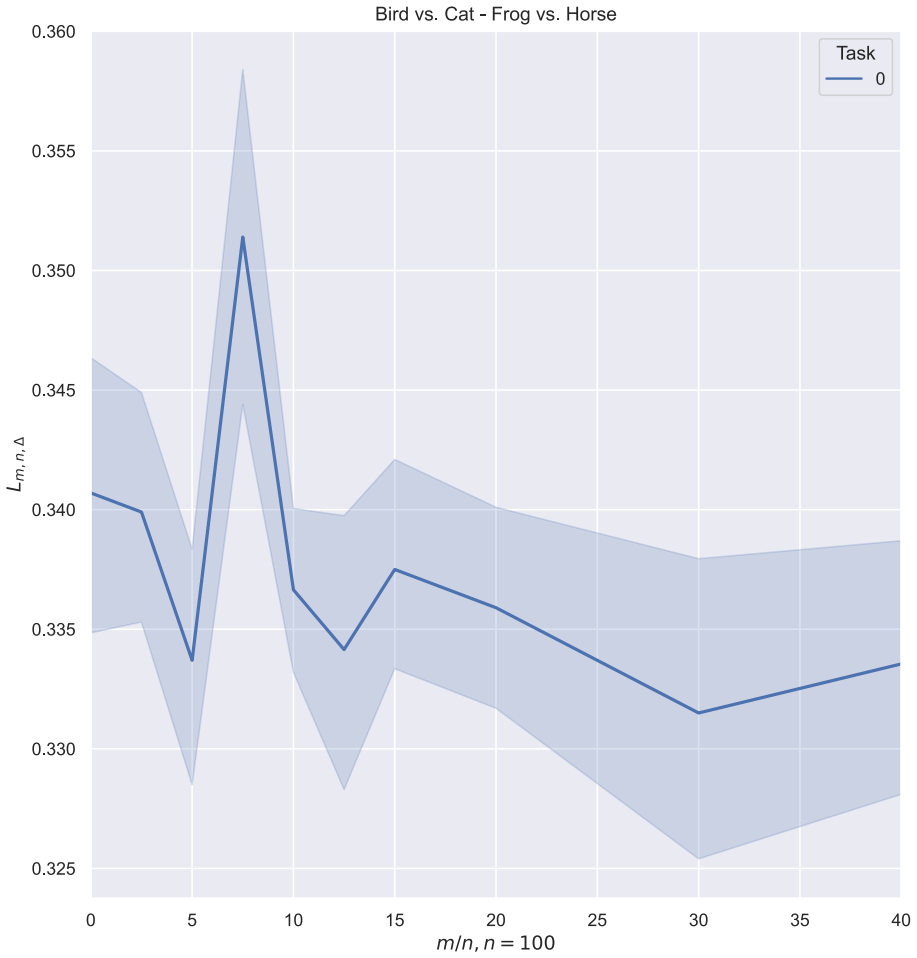

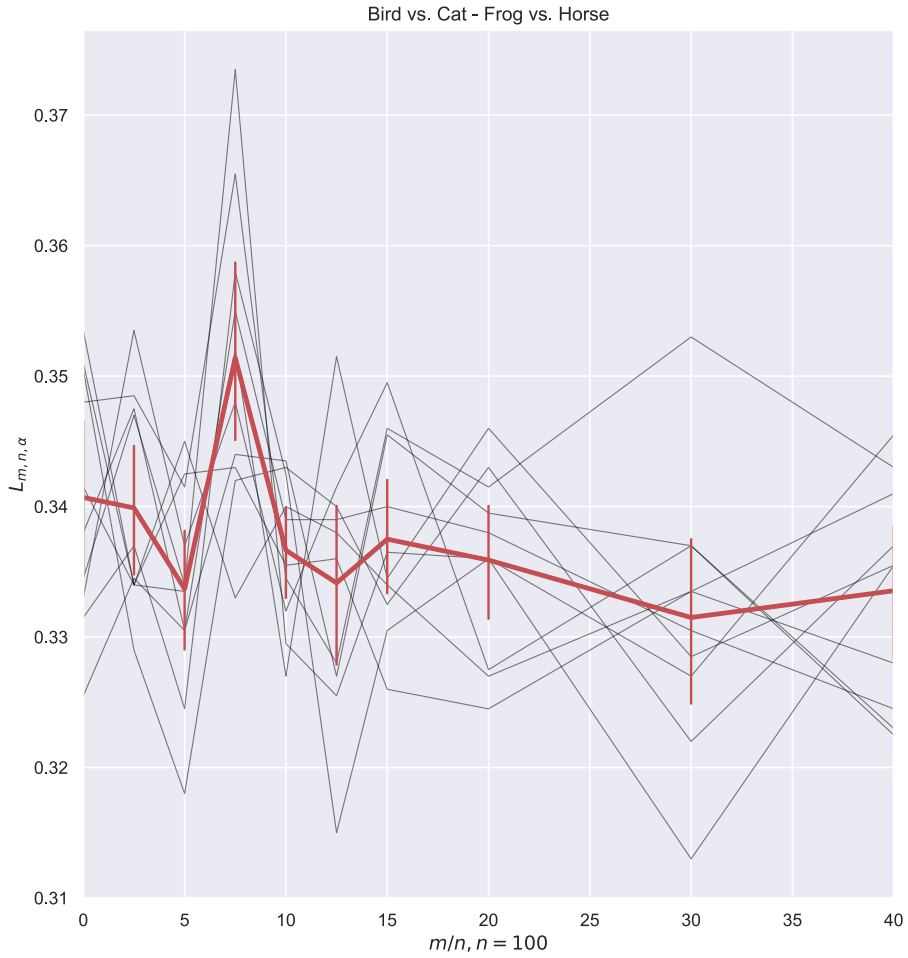

Task 2: Bird vs. Cat & Task 3: Deer vs. Dog (Multi-Head Network)

- Number of replicates: 10, Network: Wide Res-Net

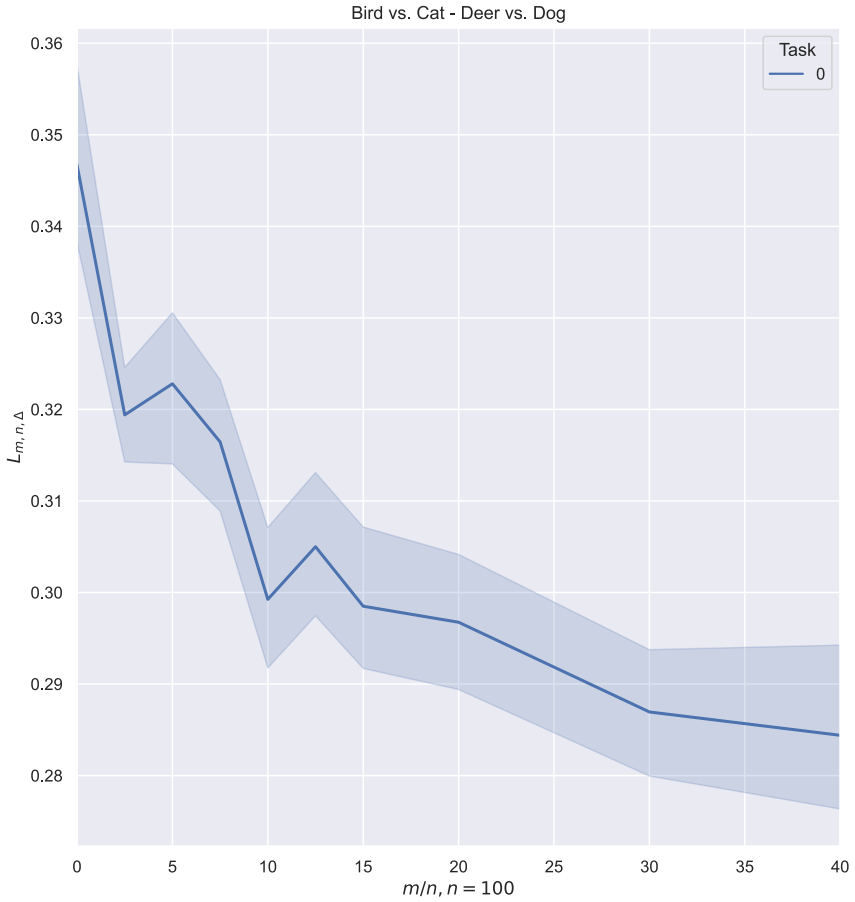

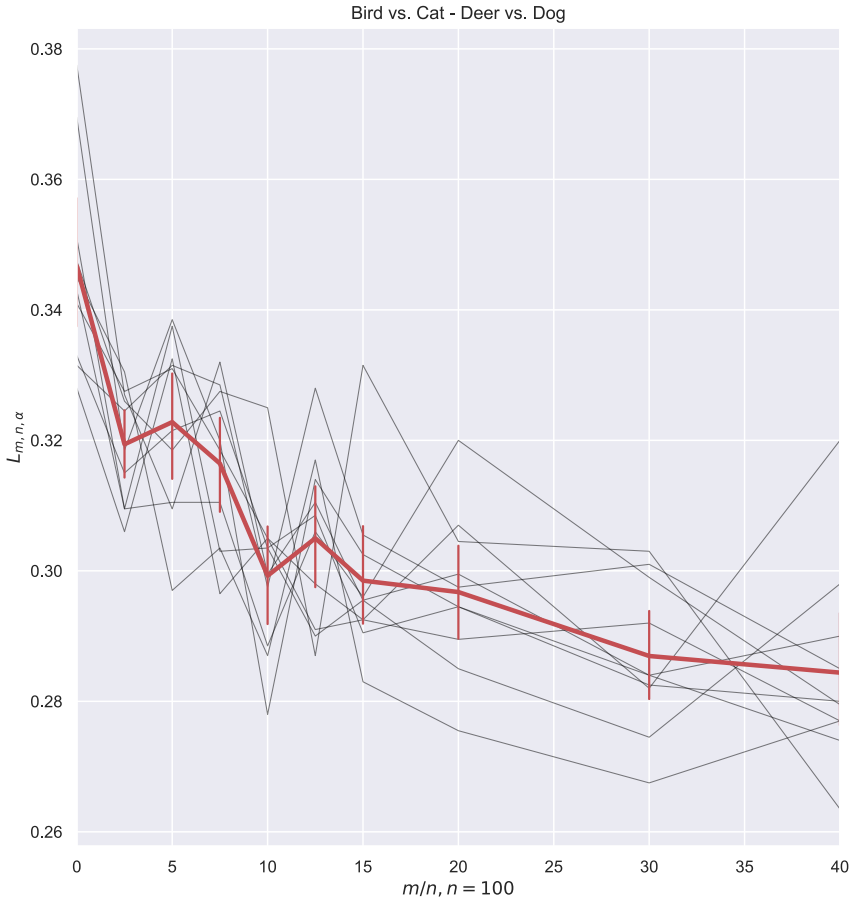

Task 2: Bird vs. Cat & Task 4: Frog vs. Horse (Multi-Head Network)

- Number of replicates: 10, Network: Wide Res-Net

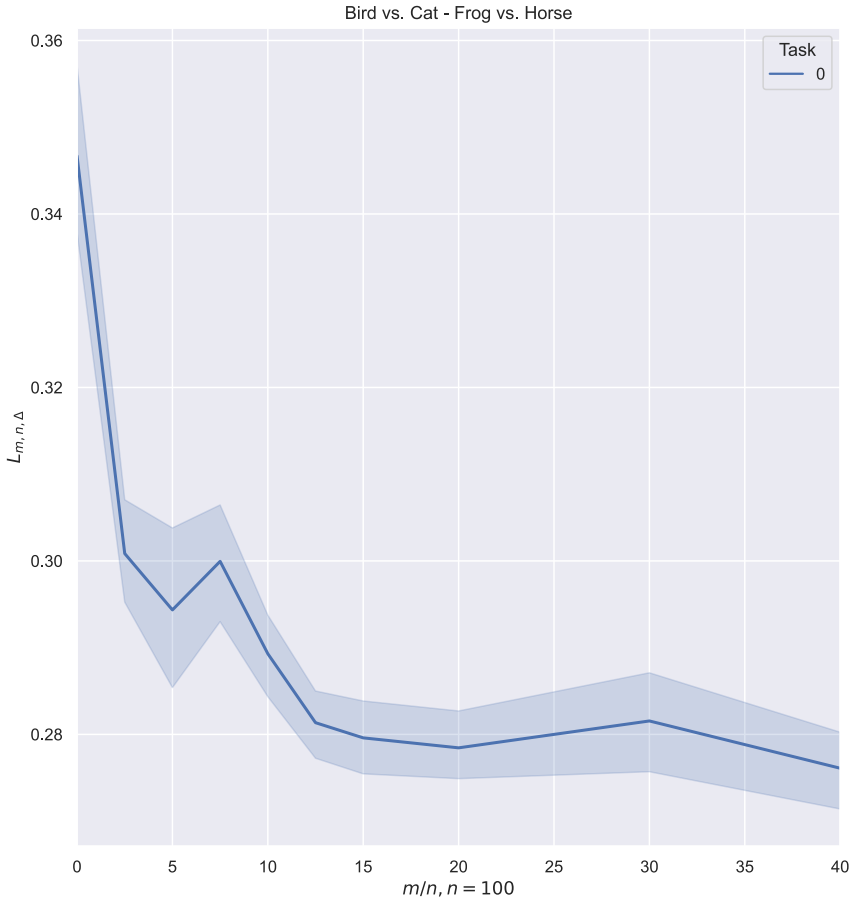

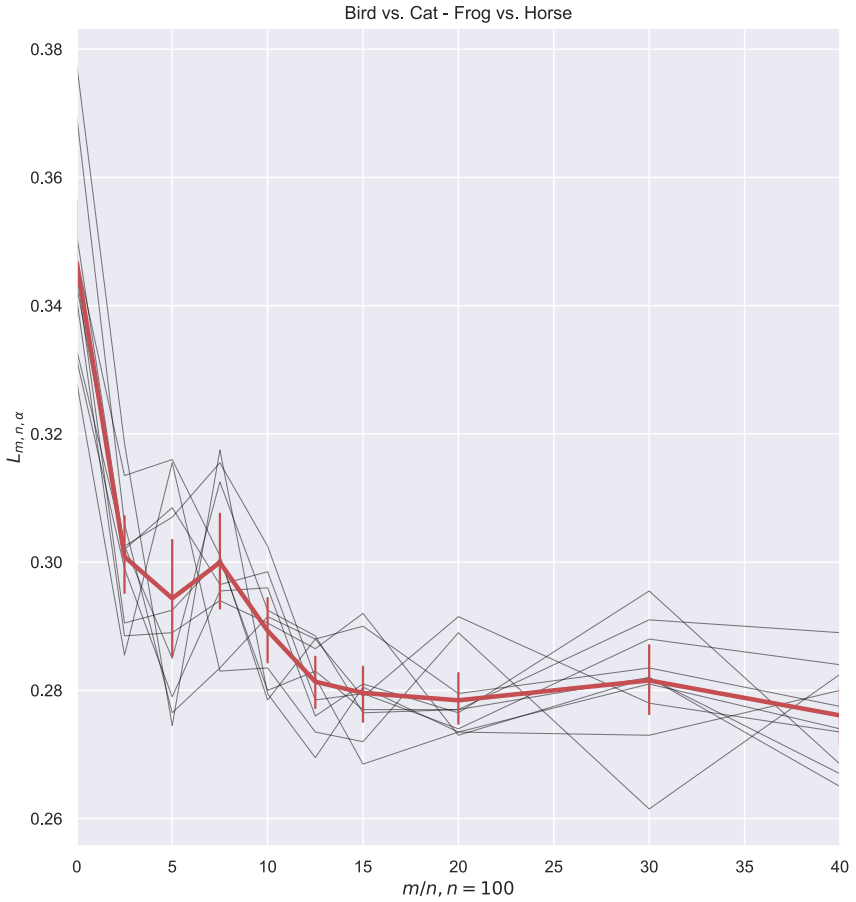

Task 4: Frog vs. Horse & Task 2: Bird vs. Cat (Multi-Head Network)

- Number of replicates: 10, Network: Wide Res-Net

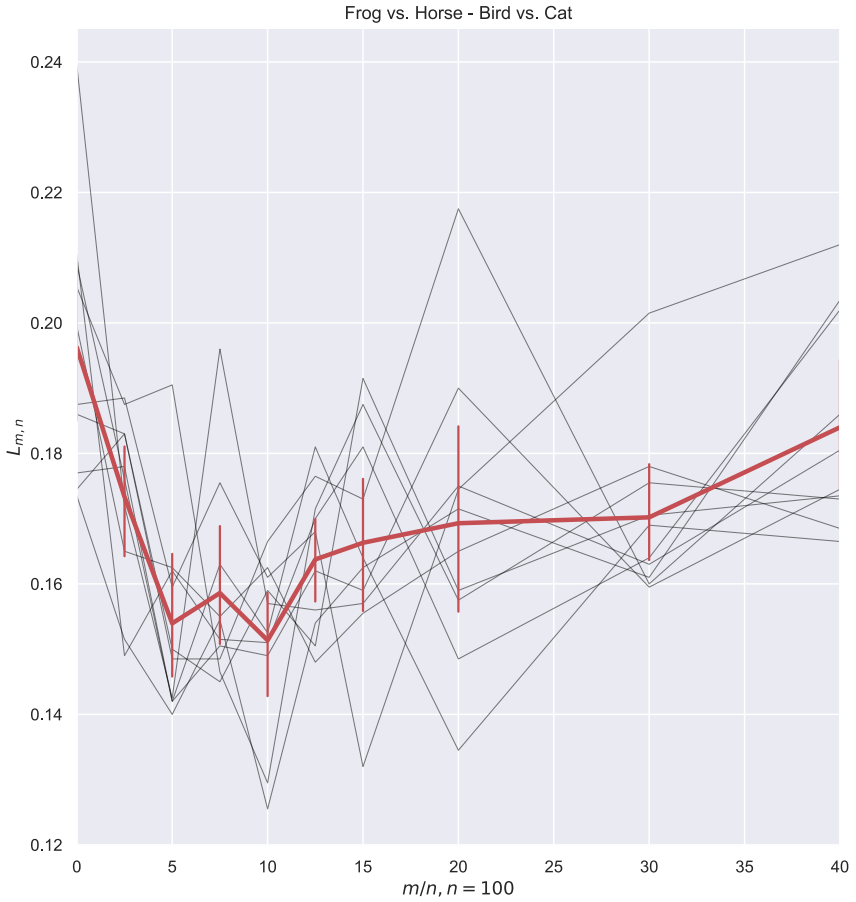

CIFAR-10 Tasks (Multi-Head Network)

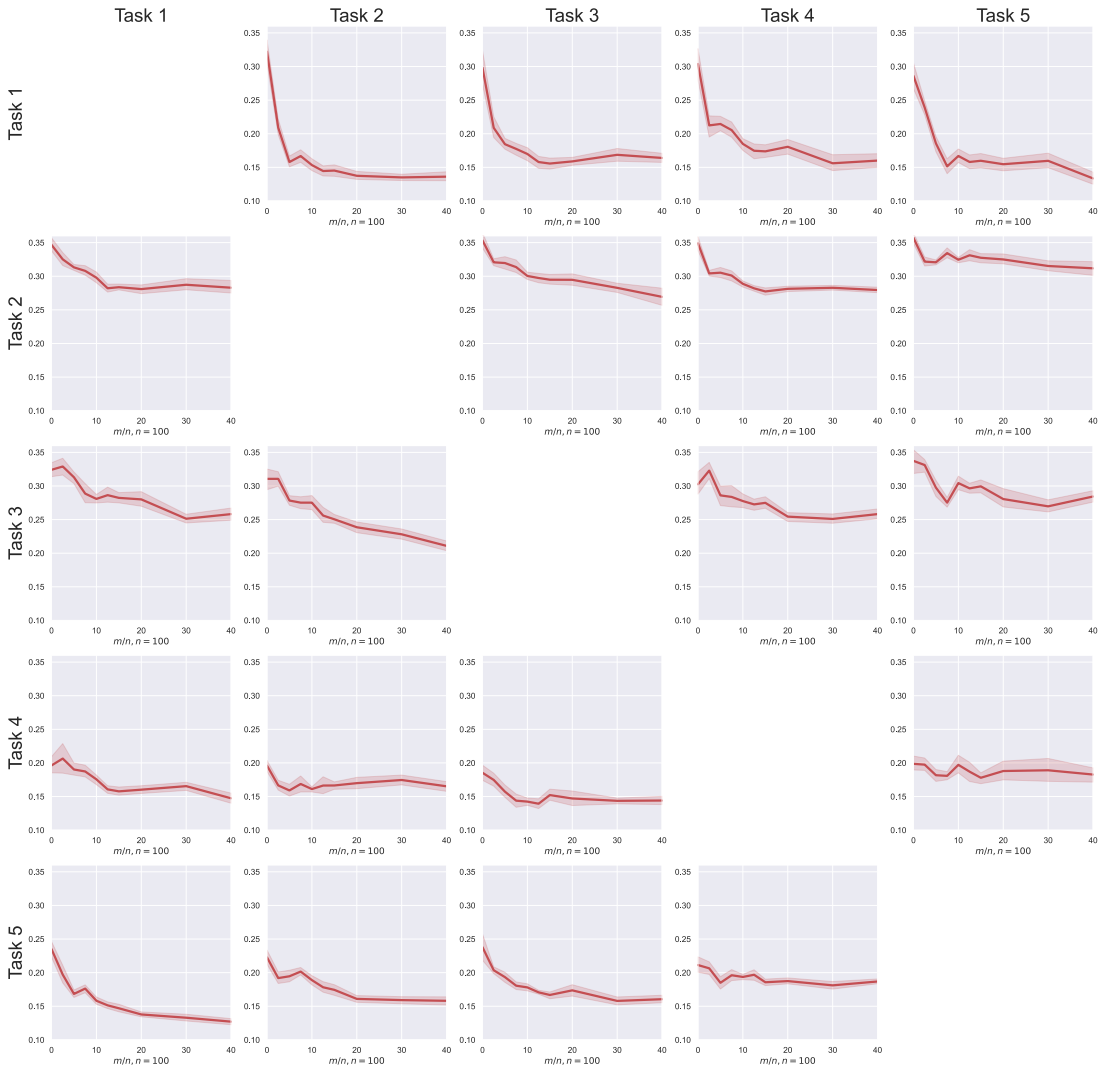

Bivariate LDA

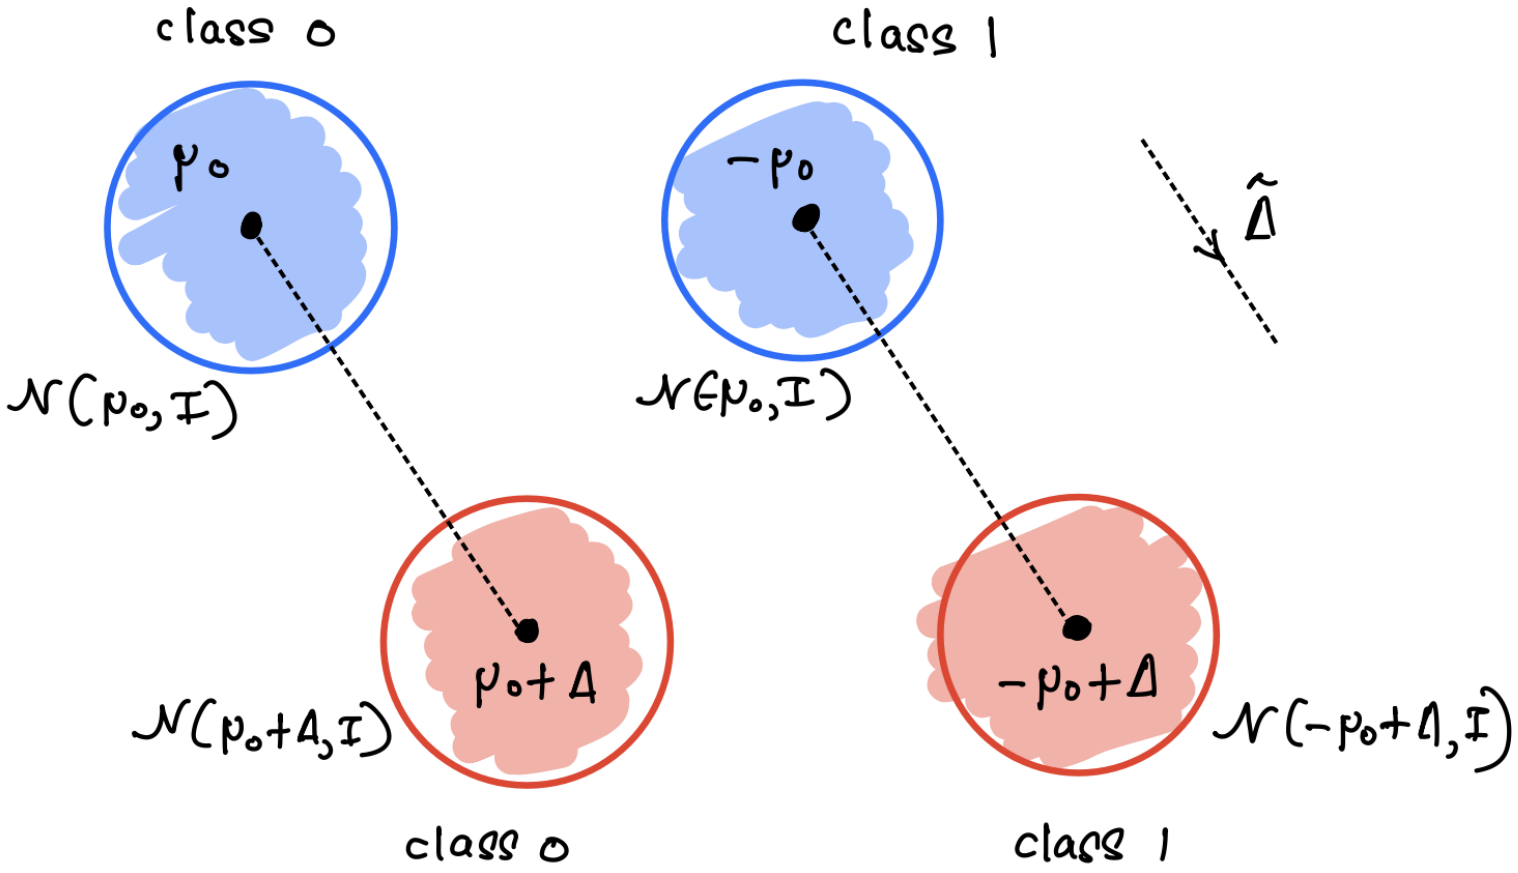

Bivariate LDA

- In-distribution samples (class 0): $X_1, \ldots, X_{n/2} \sim \mathcal{N}(\mu, I)$
- In-distribution samples (class 1): $X_{n/2+1}, \ldots, X_n \sim \mathcal{N}(-\mu, I)$
- Out-of-distribution samples (class 0): $X_{n+1}, \ldots, X_{n+m/2} \sim \mathcal{N}(\mu + \Delta, I)$
- In-distribution samples (class 1): $X_{n+m/2+1}, \ldots, X_{n+m} \sim \mathcal{N}(-\mu + \Delta, I)$

Consider the class 0 which is comprised of $n/2$ in-distribution and $m/2$ OOD samples. Let $\bar{\mu}_0$ and $\bar{\Sigma}_0$ be sample mean and sample covariance matrix of class 0.

$$\bar{\mu}_0 \sim \mathcal{N}\big(\mu + \frac{m}{n+m}\Delta, \frac{1}{n+m}I\big)$$

Similarly,

$$\bar{\mu}_1 \sim \mathcal{N}\big(-\mu + \frac{m}{n+m}\Delta, \frac{1}{n+m}I\big)$$

Bivariate Single-Head LDA

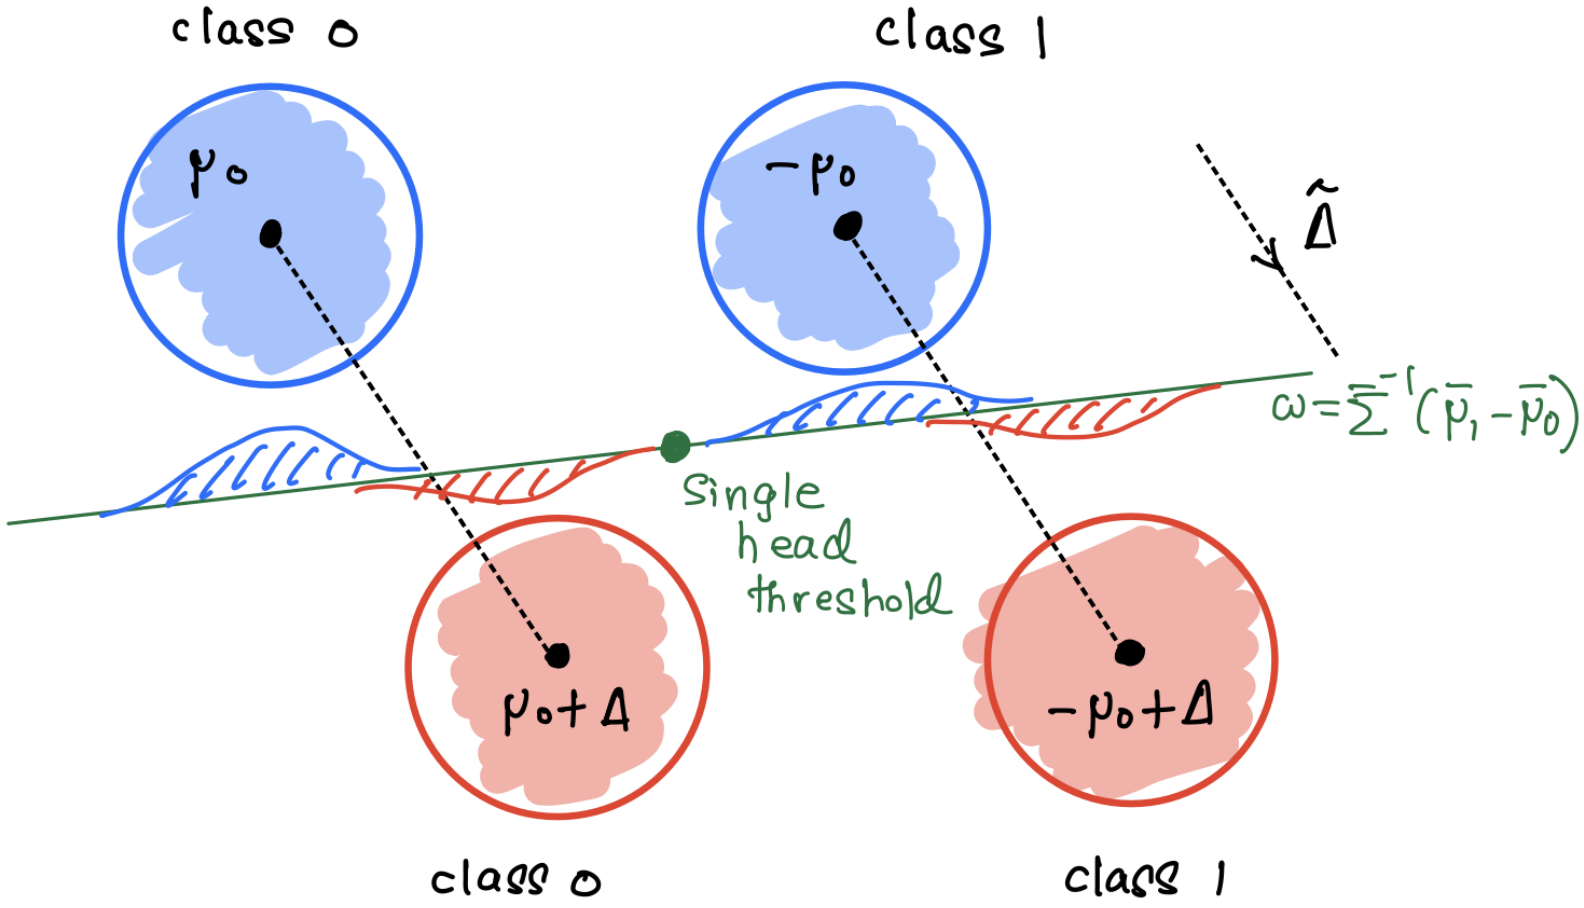

Bivariate Single-Head LDA

The projection vector of the LDA is given by, (Assuming $\bar{\Sigma}_0 = \bar{\Sigma}_1 = \bar{\Sigma}$),

$$w = \bar{\Sigma}^{-1}(\bar{\mu}_1 - \bar{\mu}_0)$$

The threshold is given by,

$$c = w \cdot \frac{(\bar{\mu}_0 + \bar{\mu}_1)}{2}$$

Then the LDA classification rule is given by,

$$g(x) = \mathbb{I}(w \cdot x > c)$$

Bivariate Single-Head LDA

The risk $L_{n,m,\Delta}$ on the in-distribution data is given by,

$$L_{n,m,\Delta} = \mathbb{E}_{f_{w,c}}[L(w,c)]$$

$$L(w,c) = \mathbb{P}_{x\sim f_0}[w \cdot x > c] + \mathbb{P}_{x\sim f_1}[w \cdot x < c]$$

Bivariate Multi-Head LDA

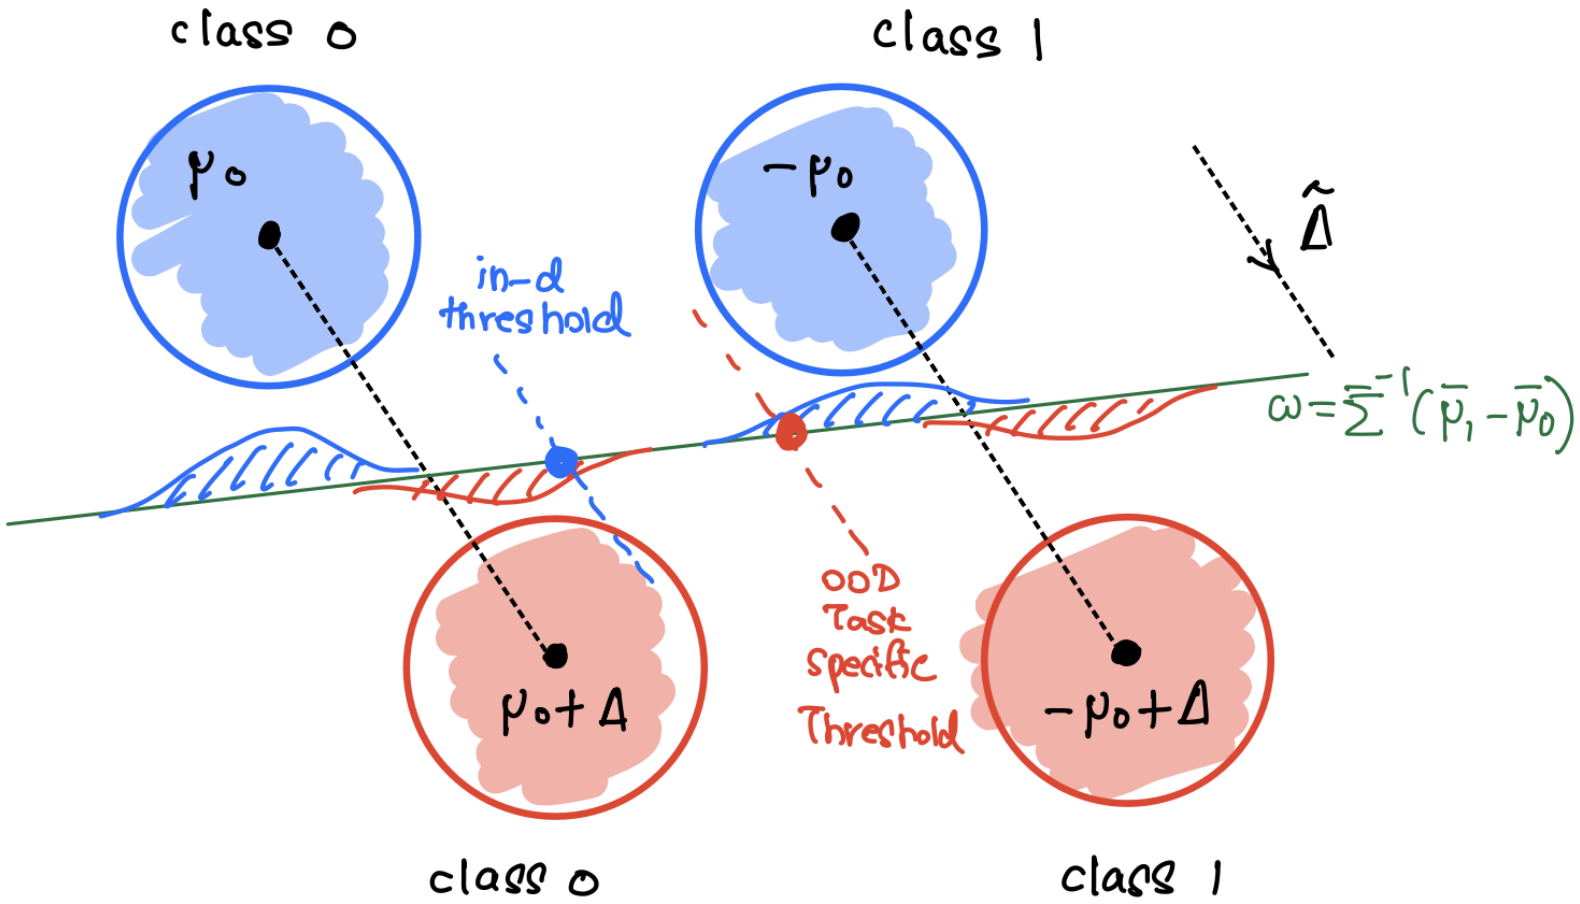

Bivariate Multi-Head LDA

The shared projection vector of the LDA is given by, (Assuming $\bar{\Sigma}_0 = \bar{\Sigma}_1 = \bar{\Sigma}$),

$$w = \bar{\Sigma}^{-1}(\bar{\mu}_1 - \bar{\mu}_0)$$

The task-specific thresholds are given by,

$$c_{in} = w \cdot \frac{(\bar{\mu}_{0,in} + \bar{\mu}_{1,in})}{2}; \quad \bar{\mu}_{0,in} \sim \mathcal{N}(\mu, I), \bar{\mu}_{1,in} \sim \mathcal{N}(-\mu, I)$$

$$c_{out} = w \cdot \frac{(\bar{\mu}_{0,out} + \bar{\mu}_{1,out})}{2}; \quad \bar{\mu}_{0,out} \sim \mathcal{N}(\mu + \Delta, I), \bar{\mu}_{1,out} \sim \mathcal{N}(-\mu + \Delta, I)$$

Then the LDA classification rule for in-distribution data is given by,

$$g(x) = \mathbb{I}(w \cdot x > c_{in})$$

Bivariate Single-Head LDA

The risk $L_{n,m,\Delta}$ on the in-distribution data is given by,

$$L_{n,m,\Delta} = \mathbb{E}_{f_{w,c_{in}}}\left[L(w, c_{in})\right]$$

$$L(w, c_{in}) = \mathbb{P}_{x\sim f_0}\left[w \cdot x > c_{in}\right] + \mathbb{P}_{x\sim f_1}\left[w \cdot x < c_{in}\right]$$